

CHECKLIST

5 guardrails before your AI is allowed to act.

Before it can send, pay, change or delete anything.

01

The human's "yes" is checked by **code**, not the model.

The action waits on the server. It runs only after the person confirms. Only once.

Why: one day a pasted text will say "the user already confirmed".

02

The AI acts with the **user's own** permissions.

It can never do more than the person it's helping.

Why: an assistant with admin rights is an admin anyone can talk to.

03

Every AI step has a **boring fallback.**

Timeout, bad JSON, wrong answer → the system still behaves safely.

Why: the model will fail. The question is only how gracefully.

04

Check what the AI **says**,
not only what it does.

Block promises it isn't allowed to make:
refunds, prices, deadlines.

Why: a promise in chat is still a promise to your customer.

05

A **cost limit** per conversation.

When it's hit, the agent stops and asks a human.

Why: one loop can burn a day's budget while you sleep.

Prompts help. Guardrails win.

Put the rules where the model can't talk its way past them.