

7 mistakes I keep seeing in **AI agent** **projects.**

And what to do instead.

01

Giving the agent **dozens of tools** “just in case”.

More tools = slower, pricier, more confused.
Give it only what this job needs.

02

Trusting **valid JSON**.

Valid format $\neq$ correct answer. Check what you DO with the answer.

03

No way to **hand off** to a human.

When the agent is stuck, it should pass the case to a person with the full context. Not loop, not guess.

04

Testing only with **tidy** examples.

Real users are messy. Test with real (anonymized) conversations.

05

Treating prompt edits as “not code”.

A prompt change is a behavior change. Test it like one.

06

No audit log.

If you can't see what the agent did and why, you can't fix it, and you can't prove it.

07

Measuring by **feeling**.

"It seems better now" is not a metric. Score it on a fixed test set.

The model is rarely the problem.

The engineering around it is.

Which one did you learn the hard way?